

# Locks and semaphores

Johan Montelius

KTH

2019

## recap, what's the problem

```

:
#include <pthread.h>

volatile int count = 0;

void *hello(void *arg) {
    for(int i = 0; i < 10; i++) {
        count++;
    }
}

int main() {
    pthread_t p1, p2;

    pthread_create(&p1, NULL, hello, NULL);
    pthread_create(&p2, NULL, hello, NULL);
    :
}
```

# Peterson's algorithm

```
int request[2] = {0,0};
int turn = 0;

int lock(int id) {

    request[id] = 1;
    int other = 1-id;
    turn = other;

    while(request[other] == 1 && turn == other) {}; // spin

    return 1;
}

void release(int id) {
    request[id] = 0;
}
```

# Total Store Order

P1



P2



# Total Store Order

P1



a  
0




0  
b




P2




# Total Store Order



# Total Store Order




# Total Store Order





# Total Store Order



# atomic memory operations

All CPU:s provide several versions of atomic operations that both read and write to a memory element in one atomic operation.

- test-and-set: swap i.e. read and write to a memory location, the simplest primitive

All CPU:s provide several versions of atomic operations that both read and write to a memory element in one atomic operation.

- test-and-set: swap i.e. read and write to a memory location, the simplest primitive
- fetch-and-add/and/xor/... : update the value with a given operation, more flexible

All CPU:s provide several versions of atomic operations that both read and write to a memory element in one atomic operation.

- test-and-set: swap i.e. read and write to a memory location, the simplest primitive
- fetch-and-add/and/xor/... : update the value with a given operation, more flexible
- compare-and-swap : if the memory location contains a specific value then swap

## try to lock by swap

```
int try(int *lock) {  
    return __sync_val_compare_and_swap(lock, 0, 1);  
}
```

## try to lock by swap

```
int try(int *lock) {  
    return __sync_val_compare_and_swap(lock, 0, 1);  
}
```

```
pushq    %rbp  
movq     %rsp, %rbp  
movq     %rdi, -8(%rbp)  
movq     -8(%rbp), %rdx  
movl     $0, %eax  
movl     $1, %ecx  
lock cmpxchgl    %ecx, (%rdx)  
nop  
popq     %rbp  
ret
```

## try to lock by swap

```
int try(int *lock) {  
    return __sync_val_compare_and_swap(lock, 0, 1);  
}
```

```
pushq    %rbp  
movq     %rsp, %rbp  
movq     %rdi, -8(%rbp)  
movq     -8(%rbp), %rdx  
movl     $0, %eax  
movl     $1, %ecx  
lock cmpxchgl    %ecx, (%rdx)  
nop  
popq     %rbp  
ret
```

*This is using GCC extensions to C, similar extensions available in all compilers.*



## a spin-lock

```
int lock(int *lock) {  
    while(try(lock) != 0) {}  
    return 1;  
}
```

## a spin-lock

```
int lock(int *lock) {  
    while(try(lock) != 0) {}  
    return 1;  
}  
  
void release(int *lock) {  
    *lock = 0;  
}
```

## finally - we're in control

```
int global = 0;

int count = 0;

void *hello(void *name) {
    for(int i = 0; i < 10; i++) {
        lock(&global);
        count++;
        release(&global);
    }
}
```





We need to talk to the operating system.

We need to talk to the operating system.

```
void lock(int *lock) {  
  
    while(try(lock) != 0) {  
        sched_yield();    // in Linux  
    }  
  
}
```

*For how long should we sleep?*



*For how long should we sleep?*



*For how long should we sleep?*



*We would like to be woken up as the lock is released - before you go-go.*



```
void lock(lock_t *m) {  
  
    while(try(m->guard) != 0) {};  
  
    if(m->flag == 0) {  
        m->flag = 1;  
        m->guard = 0;  
    } else {  
        queue_add(m->queue, gettid());  
        m->guard = 0;  
        park();  
    }  
}
```



```
void lock(lock_t *m) {  
  
    while(try(m->guard) != 0) {};  
  
    if(m->flag == 0) {  
        m->flag = 1;  
        m->guard = 0;  
    } else {  
        queue_add(m->queue, gettid());  
        m->guard = 0;  
        park();  
    }  
}
```

```
void unlock(lock_t *m) {  
  
    while(try(m->guard) != 0) {};  
  
    if(empty(m->queue)) {  
        m->flag = 0;  
    } else {  
        unpark(dequeue(m->queue));  
    }  
    m->guard = 0;  
}
```

It's not easy to to get it right.

```
/* m->flag == 1 */
:
queue_add(m->queue, gettid());
m->guard = 0;
park();
// when I wake up the flag is set

if(empty(m->queue)) {
    m->flag = 0;
} else {
    // don't reset the flag
    unpark(dequeue(m->queue));
}
```

It's not easy to to get it right.

```
/* m->flag == 1 */
:
queue_add(m->queue, gettid());
setpark();
// if someone un parks now my park() is a noop
m->guard = 0;
park();

if(empty(m->queue)) {
    m->flag = 0;
} else {
    // don't reset the flag
    unpark(dequeue(m->queue));
}
```



Introducing futex: fast user space mutex.



Introducing futex: fast user space mutex.

- `futex_wait(mutex, val)` : suspend on the `mutex` if its equal to `val`.
- `futex_wake(mutex)` : wake one of the threads suspended on the `mutex`

Introducing futex: fast user space mutex.

- `futex_wait(mutex, val)` : suspend on the `mutex` if its equal to `val`.
- `futex_wake(mutex)` : wake one of the threads suspended on the `mutex`

*In GCC you have to call them using a `syscall()`*

## a futex lock

```
void lock(volatile int *lock) {  
    while(try(lock) != 0) {  
        // time to sleep ...  
        futex_wait(lock, 1);  
    }  
}
```

## a futex lock

```
void lock(volatile int *lock) {  
    while(try(lock) != 0) {  
        // time to sleep ...  
        futex_wait(lock, 1);  
    }  
}
```

```
void unlock(volatile int *lock) {  
    *lock = 0;  
    futex_wake(lock);  
}
```

## a futex lock

```
void lock(volatile int *lock) {  
    while(try(lock) != 0) {  
        // time to sleep ...  
        futex_wait(lock, 1);  
    }  
}
```

```
void unlock(volatile int *lock) {  
    *lock = 0;  
    futex_wake(lock);  
}
```

*Not very efficient - we want to avoid calling futex\_wait() if no one is waiting.*

*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*It's better to use the pthread library API, probably more efficient and definitely less problems.*

*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*It's better to use the pthread library API, probably more efficient and definitely less problems.*

Introducing pthread mutex locks:

- `pthread_mutex_t` : structure that is the mutex



*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*It's better to use the pthread library API, probably more efficient and definitely less problems.*

Introducing pthread mutex locks:

- `pthread_mutex_t` : structure that is the mutex
- `pthread_mutex_init(pthread_mutex_t *mutex, ... *attr)`

*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*It's better to use the pthread library API, probably more efficient and definitely less problems.*

Introducing pthread mutex locks:

- `pthread_mutex_t` : structure that is the mutex
- `pthread_mutex_init(pthread_mutex_t *mutex, ... *attr)`
- `pthread_mutex_destroy(pthread_mutex_t *mutex)`

*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*It's better to use the pthread library API, probably more efficient and definitely less problems.*

Introducing pthread mutex locks:

- `pthread_mutex_t` : structure that is the mutex
- `pthread_mutex_init(pthread_mutex_t *mutex, ... *attr)`
- `pthread_mutex_destroy(pthread_mutex_t *mutex)`
- `pthread_mutex_lock(pthread_mutex_t *mutex)`

*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*It's better to use the pthread library API, probably more efficient and definitely less problems.*

Introducing pthread mutex locks:

- `pthread_mutex_t` : structure that is the mutex
- `pthread_mutex_init(pthread_mutex_t *mutex, ... *attr)`
- `pthread_mutex_destroy(pthread_mutex_t *mutex)`
- `pthread_mutex_lock(pthread_mutex_t *mutex)`
- `pthread_mutex_unlock(pthread_mutex_t *mutex)`

*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*It's better to use the pthread library API, probably more efficient and definitely less problems.*

Introducing pthread mutex locks:

- `pthread_mutex_t` : structure that is the mutex
- `pthread_mutex_init(pthread_mutex_t *mutex, ... *attr)`
- `pthread_mutex_destroy(pthread_mutex_t *mutex)`
- `pthread_mutex_lock(pthread_mutex_t *mutex)`
- `pthread_mutex_unlock(pthread_mutex_t *mutex)`

*Using Linux futex or Sun park/unpark directly is error prone and not very portable.*

*It's better to use the pthread library API, probably more efficient and definitely less problems.*

Introducing pthread mutex locks:

- `pthread_mutex_t` : structure that is the mutex
- `pthread_mutex_init(pthread_mutex_t *mutex, ... *attr)`
- `pthread_mutex_destroy(pthread_mutex_t *mutex)`
- `pthread_mutex_lock(pthread_mutex_t *mutex)`
- `pthread_mutex_unlock(pthread_mutex_t *mutex)`

*The lock procedure is platform specific, normally implemented as a combination of spinning and yield.*

# What could go wrong?

# What could go wrong?

- Nothing works, will not even compile.



# What could go wrong?

- Nothing works, will not even compile.
- Deadlock: the execution is stuck, no thread is making progress.

# What could go wrong?

- Nothing works, will not even compile.
- Deadlock: the execution is stuck, no thread is making progress.
- Livelock: we're moving around in circles, all threads think that they are doing progress but we're stuck in a loop.

# What could go wrong?

- Nothing works, will not even compile.
- Deadlock: the execution is stuck, no thread is making progress.
- Livelock: we're moving around in circles, all threads think that they are doing progress but we're stuck in a loop.
- Starvation: we're making progress but some threads are stuck waiting.

# What could go wrong?

- Nothing works, will not even compile.
- Deadlock: the execution is stuck, no thread is making progress.
- Livelock: we're moving around in circles, all threads think that they are doing progress but we're stuck in a loop.
- Starvation: we're making progress but some threads are stuck waiting.
- Unfairness: we're making progress but some threads are given more of the resources.


# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.

H: 

M:

L:




# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.

H: 


M: 

L:




# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.



# Resources, priorities and scheduling


Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.






# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.




# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.




# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.




# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.




# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.




# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.




# Resources, priorities and scheduling

Assume we have a fixed priority scheduler, three processes with high (H), medium (M) and low (L) priority and one critical resource.



# Mars Pathfinder and Priority Inversion





# Some examples

- concurrent counter
- a list
- a queue

# the concurrent counter

```
struct counter_t {  
    int val;  
}  
  
void incr(struct counter_t *c) {  
    c->val++;  
}
```

# the concurrent counter

```
struct counter_t {  
    int val;  
}
```

```
void incr(struct counter_t *c) {  
    c->val++;  
}
```

```
struct counter_t {  
    int val;  
    pthread_mutex_t lock;  
}
```


```
void incr(struct counter_t *c) {  
    pthread_lock(c->lock);  
    c->val++;  
    pthread_unlock(c->lock);  
}
```

# Do the right thing


*Doing the right thing often has a price.*

# Do the right thing


*Doing the right thing often has a price.*




# sloppy counter



# sloppy counter




# sloppy counter







# sloppy counter




# sloppy counter




# sloppy counter




# sloppy counter




# sloppy counter




# sloppy counter



# sloppy counter




# sloppy counter







# sloppy counter




# sloppy counter




# sloppy counter




# sloppy counter



# sloppy counter



# sloppy counter



*Sloppy vs Speed - do the right thing.*

*Simple solution: protect the list with one lock.*

*Simple solution: protect the list with one lock.*

*Concurrent solution: allow several thread to operate on the list concurrently.*



*Simple solution: protect the list with one lock.*

*Concurrent solution: allow several thread to operate on the list concurrently.*

- concurrent reading: not a problem

*Simple solution: protect the list with one lock.*

*Concurrent solution: allow several thread to operate on the list concurrently.*

- concurrent reading: not a problem
- concurrent updating: ....

*Simple solution: protect the list with one lock.*

*Concurrent solution: allow several thread to operate on the list concurrently.*

- concurrent reading: not a problem
- concurrent updating: .... hmm, how would you solve it?

# What about a queue


*Simple solution: protect the queue with one lock.*

*Concurrent solution: allow threads to add elements to the queue at the same time as other remove elements.*

# What about a queue

*Simple solution: protect the queue with one lock.*


*Concurrent solution: allow threads to add elements to the queue at the same time as other remove elements.*



# What about a queue

*Simple solution: protect the queue with one lock.*


*Concurrent solution: allow threads to add elements to the queue at the same time as other remove elements.*



# What about a queue

*Simple solution: protect the queue with one lock.*


*Concurrent solution: allow threads to add elements to the queue at the same time as other remove elements.*



# What about a queue

*Simple solution: protect the queue with one lock.*

*Concurrent solution: allow threads to add elements to the queue at the same time as other remove elements.*






# What about a queue

*Simple solution: protect the queue with one lock.*


*Concurrent solution: allow threads to add elements to the queue at the same time as other remove elements.*



# What about a queue

*Simple solution: protect the queue with one lock.*

*Concurrent solution: allow threads to add elements to the queue at the same time as other remove elements.*



*Traditionally operating systems were single threaded - the obvious solution.*

*Traditionally operating systems were single threaded - the obvious solution.*

*The first systems that operated on multi-cpu architectures used one **big kernel lock** to avoid any problems with concurrency.*

*Traditionally operating systems were single threaded - the obvious solution.*

*The first systems that operated on multi-cpu architectures used one **big kernel lock** to avoid any problems with concurrency.*

*An operating system that is targeting multi-core architectures will today be multi threaded and use fine grain locking to increase performance.*

*Traditionally operating systems were single threaded - the obvious solution.*

*The first systems that operated on multi-cpu architectures used one **big kernel lock** to avoid any problems with concurrency.*

*An operating system that is targeting multi-core architectures will today be multi threaded and use fine grain locking to increase performance.*

*How are things done in for example the JVM or Erlang?*

The locks that we have seen are all right:

- We can take a lock and prevent others from obtaining the lock.

The locks that we have seen are all right:

- We can take a lock and prevent others from obtaining the lock.
- If someone holds the lock we will suspend execution.



The locks that we have seen are all right:

- We can take a lock and prevent others from obtaining the lock.
- If someone holds the lock we will suspend execution.
- When the lock is released we will wake up and try to grab the lock again.


The locks that we have seen are all right:

- We can take a lock and prevent others from obtaining the lock.
- If someone holds the lock we will suspend execution.
- When the lock is released we will wake up and try to grab the lock again.


We would like to suspend and only be woken up if a specified condition holds true.




## the queue revisited




## the queue revisited



## the queue revisited



## the queue revisited



What do we do now?





Introducing pthread conditional variables:

- `pthread_cond_t` : the data structure of a conditional variable

Introducing pthread conditional variables:

- `pthread_cond_t` : the data structure of a conditional variable
- `pthread_cond_init(pthread_cond_t *restrict cond, ...)`

Introducing pthread conditional variables:

- `pthread_cond_t` : the data structure of a conditional variable
- `pthread_cond_init(pthread_cond_t *restrict cond, ...)`
- `pthread_cond_destroy(pthread_cond_t *cond)`

Introducing pthread conditional variables:

- `pthread_cond_t` : the data structure of a conditional variable
- `pthread_cond_init(pthread_cond_t *restrict cond, ...)`
- `pthread_cond_destroy(pthread_cond_t *cond)`
- `pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)`

Introducing pthread conditional variables:

- `pthread_cond_t` : the data structure of a conditional variable
- `pthread_cond_init(pthread_cond_t *restrict cond, ...)`
- `pthread_cond_destroy(pthread_cond_t *cond)`
- `pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)`
- `pthread_cond_signal(pthread_cond_t *cond)`

Introducing pthread conditional variables:

- `pthread_cond_t` : the data structure of a conditional variable
- `pthread_cond_init(pthread_cond_t *restrict cond, ...)`
- `pthread_cond_destroy(pthread_cond_t *cond)`
- `pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)`
- `pthread_cond_signal(pthread_cond_t *cond)`
- `pthread_cond_broadcast(pthread_cond_t *cond)`

Introducing pthread conditional variables:

- `pthread_cond_t` : the data structure of a conditional variable
- `pthread_cond_init(pthread_cond_t *restrict cond, ...)`
- `pthread_cond_destroy(pthread_cond_t *cond)`
- `pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)`
- `pthread_cond_signal(pthread_cond_t *cond)`
- `pthread_cond_broadcast(pthread_cond_t *cond)`

Introducing pthread conditional variables:

- `pthread_cond_t` : the data structure of a conditional variable
- `pthread_cond_init(pthread_cond_t *restrict cond, ...)`
- `pthread_cond_destroy(pthread_cond_t *cond)`
- `pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)`
- `pthread_cond_signal(pthread_cond_t *cond)`
- `pthread_cond_broadcast(pthread_cond_t *cond)`

*The exact declarations are slightly more complicated, check the man pages.*



# the producer/consumer

A single element buffer, multiple consumers, multiple producers.

```
int buffer;  
int count = 0;
```

# the producer/consumer

A single element buffer, multiple consumers, multiple producers.

```
int buffer;  
int count  = 0;
```

```
void put(int value) {  
    assert(count == 0);  
    count = 1;  
    buffer = value;  
}
```

```
int get() {  
    assert(count == 1);  
    count = 0;  
    return buffer;  
}
```

# the producer/consumer

A single element buffer, multiple consumers, multiple producers.

```
int buffer;  
int count  = 0;
```

```
void put(int value) {  
    assert(count == 0);  
    count = 1;  
    buffer = value;  
}
```

```
int get() {  
    assert(count == 1);  
    count = 0;  
    return buffer;  
}
```

*Let's try to make this work.*

this will not work

```
void produce(int val) {  
    put(val);  
}
```

```
int consume() {  
    int val = get();  
    return val;  
}
```

## add a mutex and cond variable

```
pthread_cond_t  cond;  
pthread_mutex_t mutex;
```

## add a mutex and cond variable

```
pthread_cond_t cond;  
pthread_mutex_t mutex;
```

```
produce(int val) {  
    pthread_mutex_lock(&mutex);  
    if(count == 1)  
        pthread_cond_wait(&cond, &mutex);  
    put(val);  
    pthread_cond_signal(&cond);  
    pthread_mutex_unlock(&mutex);  
}
```

## add a mutex and cond variable

```
pthread_cond_t cond;  
pthread_mutex_t mutex;
```

```
produce(int val) {  
    pthread_mutex_lock(&mutex);  
    if(count == 1)  
        pthread_cond_wait(&cond, &mutex);  
    put(val);  
    pthread_cond_signal(&cond);  
    pthread_mutex_unlock(&mutex);  
}
```

```
int consume() {  
    pthread_mutex_lock(&mutex);  
    if(count == 0)  
        pthread_cond_wait(&cond, &mutex);  
    int val = get();  
    pthread_cond_signal(&cond);  
    pthread_mutex_unlock(&mutex);  
    return val;  
}
```

## add a mutex and cond variable

```
pthread_cond_t cond;  
pthread_mutex_t mutex;
```

```
produce(int val) {  
    pthread_mutex_lock(&mutex);  
    if(count == 1)  
        pthread_cond_wait(&cond, &mutex);  
    put(val);  
    pthread_cond_signal(&cond);  
    pthread_mutex_unlock(&mutex);  
}
```

```
int consume() {  
    pthread_mutex_lock(&mutex);  
    if(count == 0)  
        pthread_cond_wait(&cond, &mutex);  
    int val = get();  
    pthread_cond_signal(&cond);  
    pthread_mutex_unlock(&mutex);  
    return val;  
}
```

*When does this work, when does it not work?*



## a race condition

If you're signaled to wake up - it might take some time before you do wake up.



```
pthread_cond_t filled, empty;  
pthread_mutex_t mutex;
```

```
pthread_cond_t filled, empty;  
pthread_mutex_t mutex;
```

```
produce(int val) {  
    pthread_mutex_lock(&mutex);  
    while(count == 1)  
        pthread_cond_wait(&empty, &mutex);  
    :  
    pthread_cond_signal(&filled);  
    :  
}
```

```
pthread_cond_t filled, empty;  
pthread_mutex_t mutex;
```

```
produce(int val) {  
    pthread_mutex_lock(&mutex);  
    while(count == 1)  
        pthread_cond_wait(&empty, &mutex);  
    :  
    pthread_cond_signal(&filled);  
    :  
}
```

```
int consume() {  
    pthread_mutex_lock(&mutex);  
    while(count == 0)  
        pthread_cond_wait(&filled, &mutex);  
    :  
    pthread_cond_signal(&empty);  
    :  
}
```

## a larger buffer

```
int buffer[MAX];  
int *getp = 0;  
int *putp = 0;  
int count = 0;
```

```
void put(int value) {  
    assert(count < MAX);  
    buffer[putp] = value;  
    putp = putp + 1 % MAX;  
    count++;  
}
```

```
int get() {  
    assert(count > 0);  
    int val = buffer[getp];  
    getp = getp + 1 % MAX;  
    count--;  
    return val;  
}
```

```
produce(int val) {  
    :  
    while(count == MAX)  
        pthread_cond_wait(&empty, &mutex);  
    :  
}
```

```
produce(int val) {  
    :  
    while(count == MAX)  
        pthread_cond_wait(&empty, &mutex);  
    :  
}
```

```
int consume() {  
    :  
    while(count == 0)  
        pthread_cond_wait(&filled, &mutex);  
    :  
}
```

```
produce(int val) {  
    :  
    while(count == MAX)  
        pthread_cond_wait(&empty, &mutex);  
    :  
}
```

```
int consume() {  
    :  
    while(count == 0)  
        pthread_cond_wait(&filled, &mutex);  
    :  
}
```

*Can we allow a producer to add an entry while another removes an entry?*



# Where are we now?

- atomic test and set: we need it

# Where are we now?

- atomic test and set: we need it
- spin locks: simple to use but have some problems

# Where are we now?

- atomic test and set: we need it
- spin locks: simple to use but have some problems
- wait and wake : avoid spinning

# Where are we now?


- atomic test and set: we need it
- spin locks: simple to use but have some problems
- wait and wake : avoid spinning
- condition variables : don't wake up if it's not time to continue


# Where are we now?

- atomic test and set: we need it
- spin locks: simple to use but have some problems
- wait and wake : avoid spinning
- condition variables : don't wake up if it's not time to continue


*Is there more?*

# Semaphores






Properties of a semaphore:



Properties of a semaphore:


- holds a number






Properties of a semaphore:

- holds a number
- only allow threads to pass is number is above 0




Properties of a semaphore:

- holds a number
- only allow threads to pass is number is above 0
- passing threads decremented the number



Properties of a semaphore:

- holds a number
- only allow threads to pass is number is above 0
- passing threads decremented the number
- a thread can increment the number



Properties of a semaphore:

- holds a number
- only allow threads to pass is number is above 0
- passing threads decremented the number
- a thread can increment the number

*A semaphore is a counter of resources.*

- `#include <semaphore.h>`

# POSIX semaphores

- `#include <semaphore.h>`
- `sem_t` : the semaphore data structure

- `#include <semaphore.h>`
- `sem_t` : the semaphore data structure
- `sem_init(sem_t *sem, int pshared, unsigned int value)`: could be shared between processes

- `#include <semaphore.h>`
- `sem_t` : the semaphore data structure
- `sem_init(sem_t *sem, int pshared, unsigned int value)`: could be shared between processes
- `int sem_destroy(sem_t *sem)`



- `#include <semaphore.h>`
- `sem_t` : the semaphore data structure
- `sem_init(sem_t *sem, int pshared, unsigned int value)`: could be shared between processes
- `int sem_destroy(sem_t *sem)`
- `sem_wait(sem_t *sem)`

- `#include <semaphore.h>`
- `sem_t` : the semaphore data structure
- `sem_init(sem_t *sem, int pshared, unsigned int value)`: could be shared between processes
- `int sem_destroy(sem_t *sem)`
- `sem_wait(sem_t *sem)`
- `sem_post(sem_t *sem)`

- `#include <semaphore.h>`
- `sem_t` : the semaphore data structure
- `sem_init(sem_t *sem, int pshared, unsigned int value)`: could be shared between processes
- `int sem_destroy(sem_t *sem)`
- `sem_wait(sem_t *sem)`
- `sem_post(sem_t *sem)`



# Summary




# Summary



# Summary



# Summary





# Summary

