

Storage: HDD, SSD and RAID

Johan Montelius

KTH

2019


Why?

Why?


Give me two reasons why we would like to have secondary storage?


Computer architecture

Gigabyte Z170 Gaming




- 2 PCIe x16/x4
- 4 PCIe x1
- 2 USB 3.1
- 6 USB 3.0
- 4 USB 2.0
- 6 SATA-III
- 2 SATA Express
- 1 M.2
- 1 gigabit Ethernet
- 4 DDR4 SDRAM






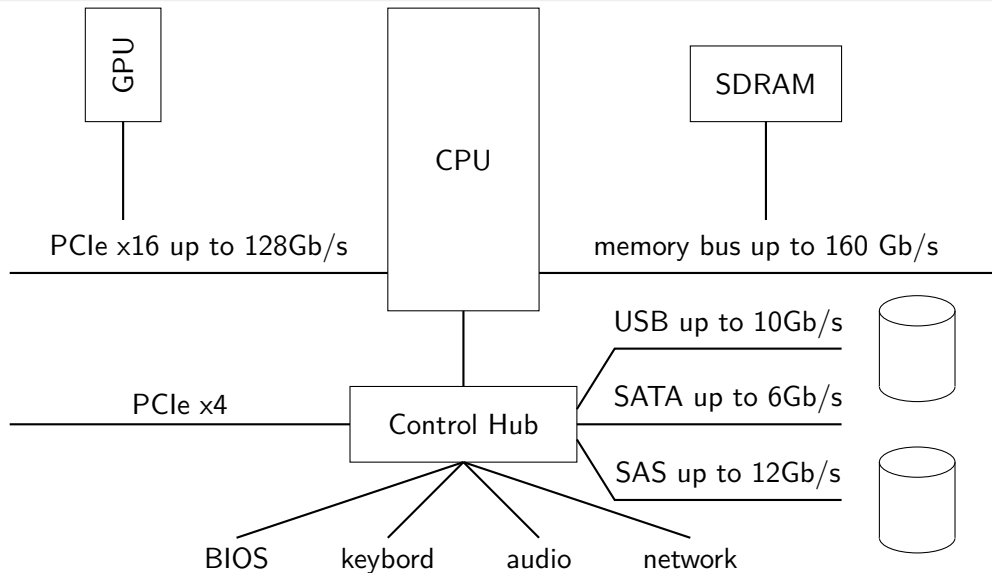
Computer architecture



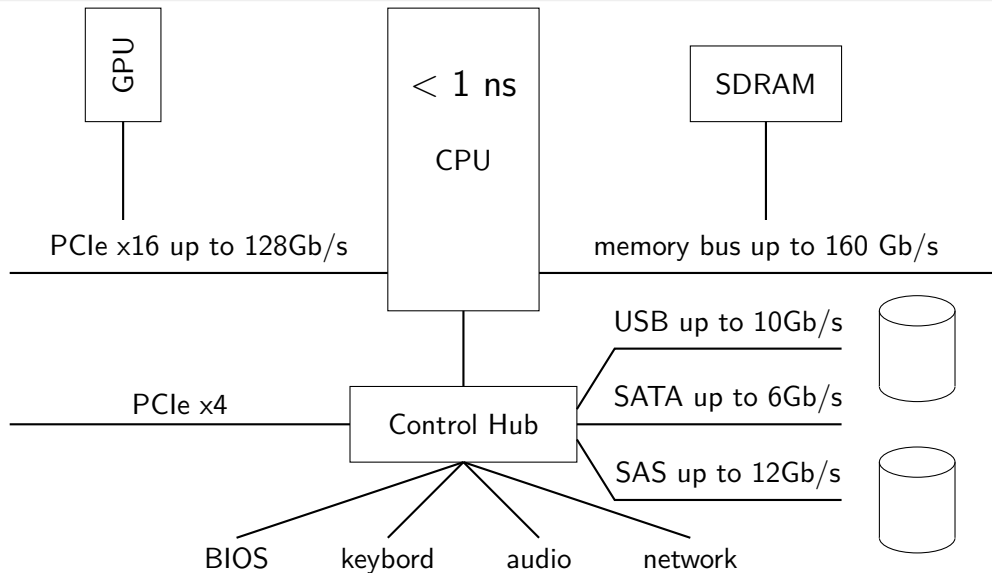
Computer architecture



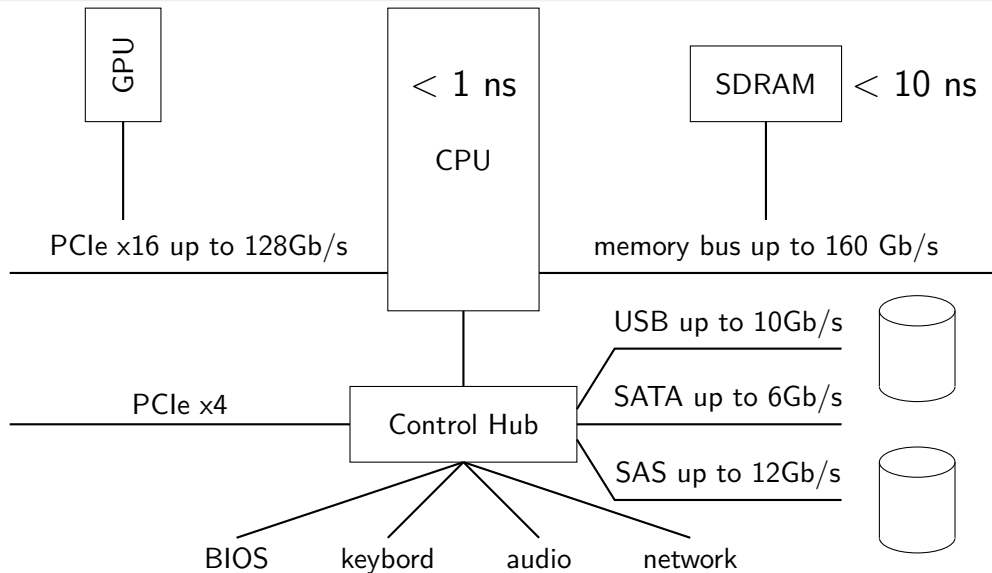
Computer architecture



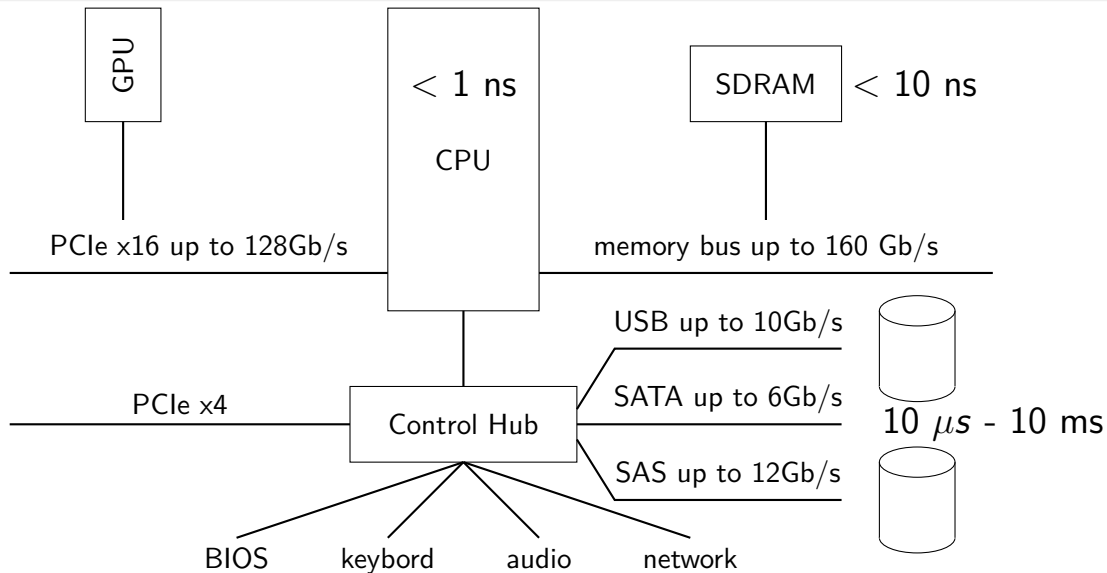
Computer architecture




Computer architecture




Computer architecture






70 percent of the code of an operating system is code for device drivers.

System architecture




70 percent of the code of an operating system is code for device drivers.

System architecture



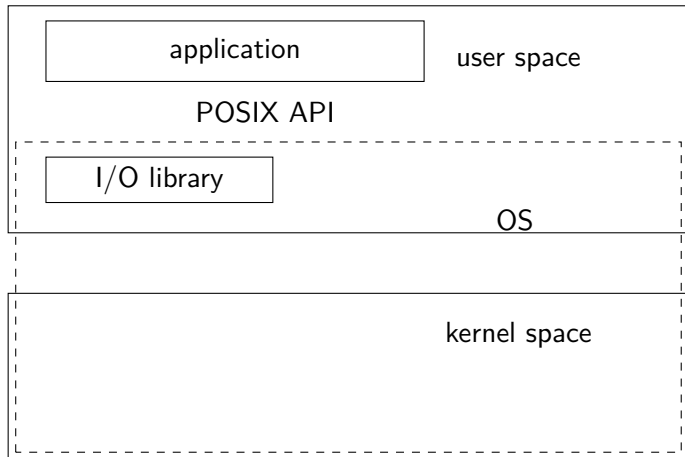
70 percent of the code of an operating system is code for device drivers.

System architecture




70 percent of the code of an operating system is code for device drivers.

System architecture




70 percent of the code of an operating system is code for device drivers.

System architecture



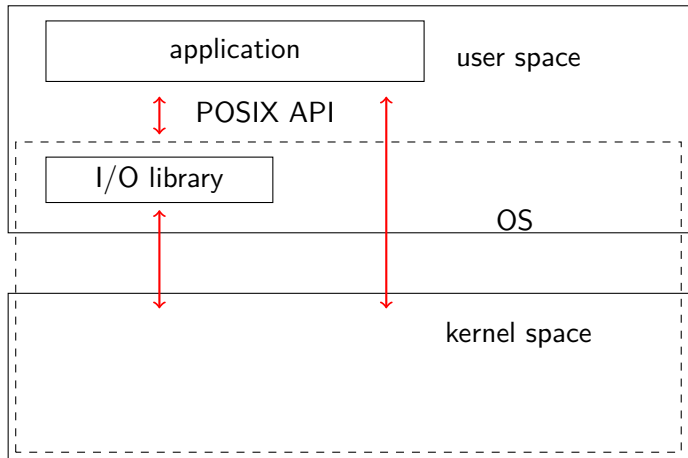
70 percent of the code of an operating system is code for device drivers.

System architecture



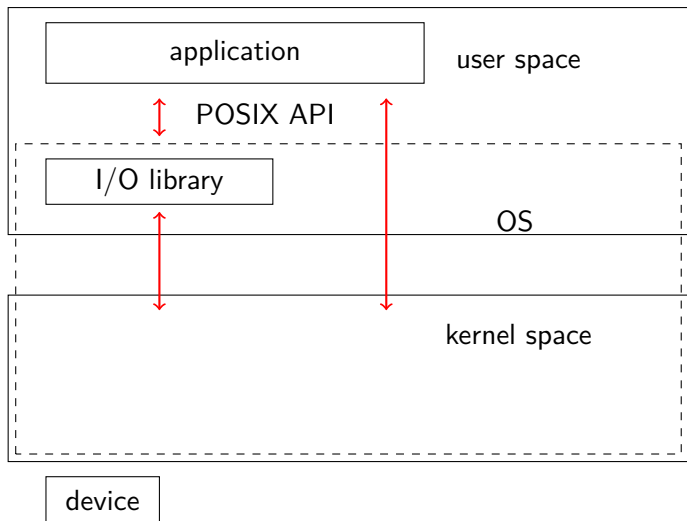
70 percent of the code of an operating system is code for device drivers.

System architecture



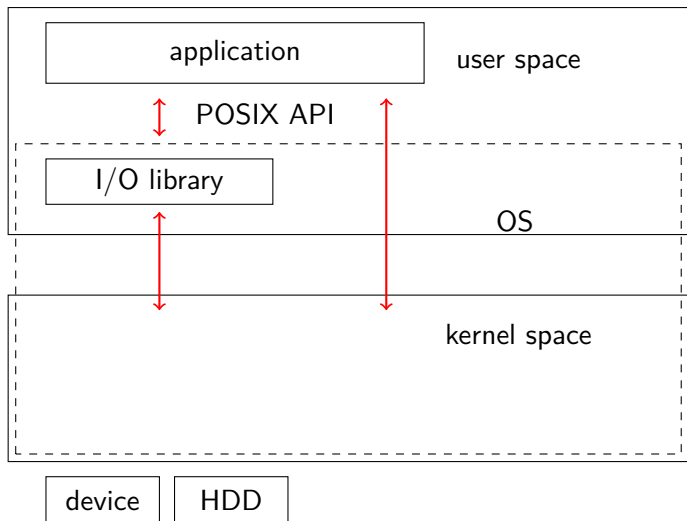
70 percent of the code of an operating system is code for device drivers.

System architecture



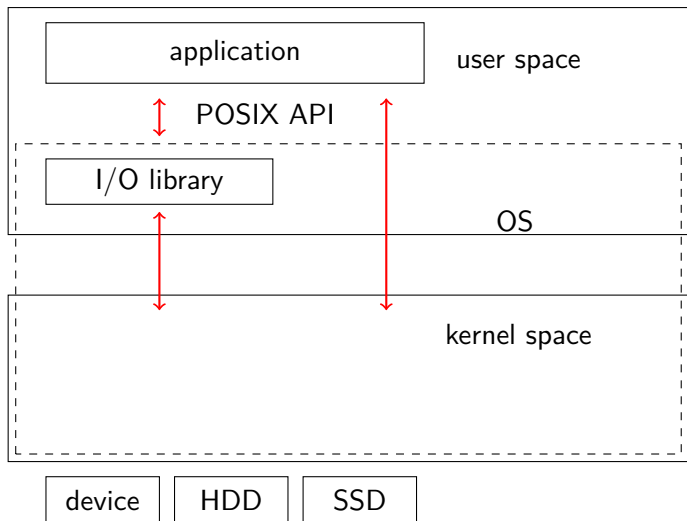
70 percent of the code of an operating system is code for device drivers.

System architecture



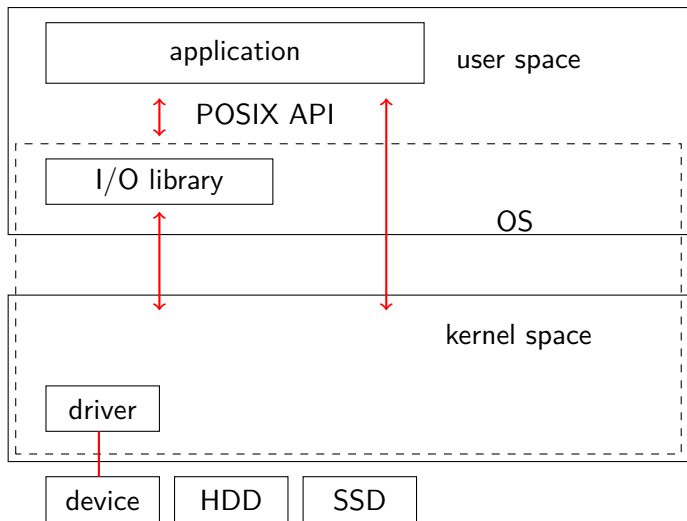
70 percent of the code of an operating system is code for device drivers.

System architecture



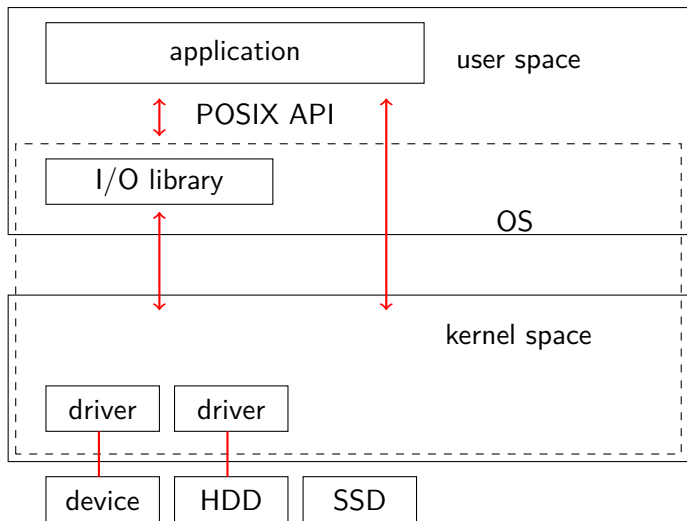
70 percent of the code of an operating system is code for device drivers.

System architecture



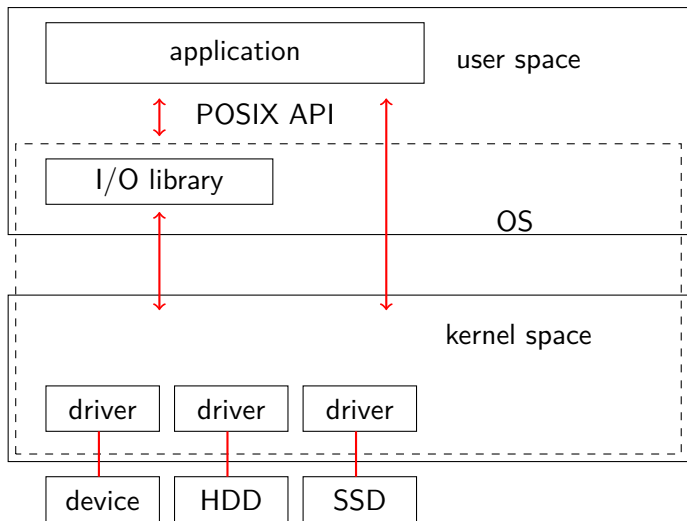
70 percent of the code of an operating system is code for device drivers.

System architecture



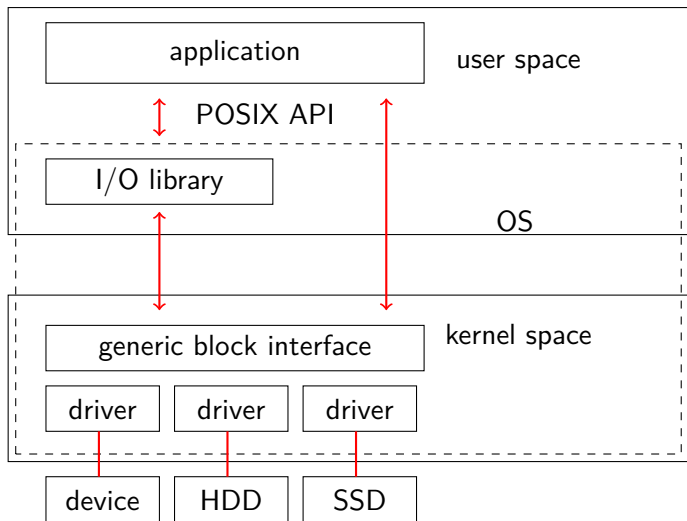
70 percent of the code of an operating system is code for device drivers.

System architecture



70 percent of the code of an operating system is code for device drivers.

System architecture




70 percent of the code of an operating system is code for device drivers.

how to interact with a device

driver

device

how to interact with a device



- A register to read the status of the device.



how to interact with a device

```
graph TD; driver[driver] --- device[device]; subgraph device; status[status]; command[command]; end
```

driver


- A register to read the status of the device.
- A register to instruct the device to read or write.

status


command

device


how to interact with a device



- A register to read the status of the device.
- A register to instruct the device to read or write.
- A register that holds the data.




how to interact with a device



- A register to read the status of the device.
- A register to instruct the device to read or write.
- A register that holds the data.
- I/O-bus could be separate from memory bus (or the same).

how to interact with a device



- A register to read the status of the device.
- A register to instruct the device to read or write.
- A register that holds the data.
- I/O-bus could be separate from memory bus (or the same).
- The driver will use either special I/O instructions or regular load/store instructions.

if you have the time

```
char read_from_device() {  
  
    while(STATUS == BUSY) {} // do nothing, just wait  
  
    COMMAND = READ;  
  
    while(STATUS == BUSY) {} // do nothing, just wait  
  
    return DATA;  
  
}
```

```
int read_request(int pid, char *buffer) {  
  
    while (STATUS == BUSY) {}  
  
    COMMAND = READ;  
  
    interrupt->process = pid;  
    interrupt->buffer = buffer;  
  
    block_process(pid);  
  
    scheduler();  
}
```


asynchronous I/O and interrupts


```
int interrupt_handler() {  
  
    int pid = interrupt->pid;  
    *(interrupt->buffer) = DATA;  
  
    ready_process(pid);  
}
```

asynchronous I/O and interrupts

```
int interrupt_handler() {  
  
    int pid = interrupt->pid;  
    *(interrupt->buffer) = DATA;  
  
    ready_process(pid);  
}
```

This is very schematic, more complicated in real life.





The kernel is interrupt driven.

Direct Memory Access

Allow devices to read and write to buffers in physical memory.

Direct Memory Access

Allow devices to read and write to buffers in physical memory.

```
int write_request(int pid, char *string, int size) {  
    while (STATUS == BUSY) {}  
  
    memcpy(string, buffer, size)  
  
    COMMAND = WRITE;  
  
    blocked->pid = pid;  
  
    block_process(pid);  
  
    scheduler();  
}
```

Direct Memory Access

Allow devices to read and write to buffers in physical memory.

```
int write_request(int pid, char *string, int size) {  
    while (STATUS == BUSY) {}  
  
    memcpy(string, buffer, size)  
  
    COMMAND = WRITE;  
  
    blocked->pid = pid;  
  
    block_process(pid);  
  
    scheduler();  
}
```

Each physical device is controlled by a *device driver* that provides the abstraction of a *character device* or *block device*.

Each physical device is controlled by a *device driver* that provides the abstraction of a *character device* or *block device*.

Block devices used as interface to disk drives that provide persistent storage.

Each physical device is controlled by a *device driver* that provides the abstraction of a *character device* or *block device*.

Block devices used as interface to disk drives that provide persistent storage.

All though all storage devices are presented using the same abstraction, they have very different characteristics.


Each physical device is controlled by a *device driver* that provides the abstraction of a *character device* or *block device*.

Block devices used as interface to disk drives that provide persistent storage.


All though all storage devices are presented using the same abstraction, they have very different characteristics.

To understand the challenges and options of the operating system, you should know the basics of how storage devices work.


Anatomy of a HDD




Anatomy of a HDD




Anatomy of a HDD




Anatomy of a HDD



Anatomy of a HDD




Anatomy of a HDD




- track/cylinder

Anatomy of a HDD




- track/cylinder
- sectors per track varies

Anatomy of a HDD




- track/cylinder
- sectors per track varies
- sector size: 4K or 512 bytes

Anatomy of a HDD




- track/cylinder
- sectors per track varies
- sector size: 4K or 512 bytes
- platters: 1 to 6
- heads: one side or two sides

Anatomy of a HDD



- track/cylinder
- sectors per track varies
- sector size: 4K or 512 bytes
- platters: 1 to 6
- heads: one side or two sides

Anatomy of a HDD



- track/cylinder
- sectors per track varies
- sector size: 4K or 512 bytes
- platters: 1 to 6
- heads: one side or two sides

Only one head at a time is used (no parallel read).

Sector addressing

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:

Sector addressing

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)

Sector addressing

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)

Sector addressing

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)
 - sectors per cylinder: 63 (6-bits)

Sector addressing

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)
 - sectors per cylinder: 63 (6-bits)
 - number of sectors: 1 Mi

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)
 - sectors per cylinder: 63 (6-bits)
 - number of sectors: 1 Mi
 - largest disk assuming 512 Byte sectors: 512 MiByte

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)
 - sectors per cylinder: 63 (6-bits)
 - number of sectors: 1 Mi
 - largest disk assuming 512 Byte sectors: 512 MiByte

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)
 - sectors per cylinder: 63 (6-bits)
 - number of sectors: 1 Mi
 - largest disk assuming 512 Byte sectors: 512 MiByte
- Today, sectors are addresses linearly 0.. n, Linear Block Addressing (LBA):

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)
 - sectors per cylinder: 63 (6-bits)
 - number of sectors: 1 Mi
 - largest disk assuming 512 Byte sectors: 512 MiByte
- Today, sectors are addresses linearly 0.. n, Linear Block Addressing (LBA):
 - 28-bit or 48-bit address

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)
 - sectors per cylinder: 63 (6-bits)
 - number of sectors: 1 Mi
 - largest disk assuming 512 Byte sectors: 512 MiByte
- Today, sectors are addresses linearly 0.. n, Linear Block Addressing (LBA):
 - 28-bit or 48-bit address
 - up to 256 Ti sectors

Sector addressing

- Historically sectors address by cylinder-head-sector (CHS), due to incompatible standards the limitation was:
 - cylinder: 1024 (10-bits)
 - heads: 16 (4-bits)
 - sectors per cylinder: 63 (6-bits)
 - number of sectors: 1 Mi
 - largest disk assuming 512 Byte sectors: 512 MiByte
- Today, sectors are addresses linearly 0.. n, Linear Block Addressing (LBA):
 - 28-bit or 48-bit address
 - up to 256 Ti sectors
 - largest disk assuming 4 KiByte sectors: 1 PiByte

```
> sudo hdparm -I /dev/sda  
> dmesg | grep ata2
```

HDD - Hard Disk Drive

Seagate Desktop



HDD - Hard Disk Drive

Seagate Desktop



- total capacity: 2 TiByte

HDD - Hard Disk Drive

Seagate Desktop



- total capacity: 2 TiByte
- form factor: 3.5"

HDD - Hard Disk Drive

Seagate Desktop



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm

HDD - Hard Disk Drive

Seagate Desktop



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA III

HDD - Hard Disk Drive

Seagate Desktop



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA III
- cache size: 64 MiByte

HDD - Hard Disk Drive

Seagate Desktop



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA III
- cache size: 64 MiByte
- read throughput: 156 MByte/s

HDD - Hard Disk Drive

Seagate Desktop



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA III
- cache size: 64 MiByte
- read throughput: 156 MByte/s

aprx price, October 2016, 900:-

HDD - Hard Disk Drive

Seagate Cheetah 15K



HDD - Hard Disk Drive

Seagate Cheetah 15K



- total capacity: 600 GiByte

HDD - Hard Disk Drive

Seagate Cheetah 15K



- total capacity: 600 GiByte
- form factor: 3.5"

HDD - Hard Disk Drive

Seagate Cheetah 15K



- total capacity: 600 GiByte
- form factor: 3.5"
- rotational speed: 15.000 rpm

HDD - Hard Disk Drive

Seagate Cheetah 15K



- total capacity: 600 GiByte
- form factor: 3.5"
- rotational speed: 15.000 rpm
- connection: SAS-3

HDD - Hard Disk Drive

Seagate Cheetah 15K



- total capacity: 600 GiByte
- form factor: 3.5"
- rotational speed: 15.000 rpm
- connection: SAS-3
- cache size: 16 MiByte

HDD - Hard Disk Drive

Seagate Cheetah 15K



- total capacity: 600 GiByte
- form factor: 3.5"
- rotational speed: 15.000 rpm
- connection: SAS-3
- cache size: 16 MiByte
- read throughput: 204 MByte/s


HDD - Hard Disk Drive


Seagate Cheetah 15K




- total capacity: 600 GiByte
- form factor: 3.5"
- rotational speed: 15.000 rpm
- connection: SAS-3
- cache size: 16 MiByte
- read throughput: 204 MByte/s


aprx price, October 2016, 2.200:-








- seek time: time to move arm to the right cylinder



- seek time: time to move arm to the right cylinder
- rotation time: time to rotate the disk



- seek time: time to move arm to the right cylinder
- rotation time: time to rotate the disk
- read time: read one or more sectors

- Seagate Desktop
- Seagate Cheeta 15K

HDD - shoot out

- Seagate Desktop
- rotation speed: 7200 rpm
- Seagate Cheeta 15K
- rotation speed: 15000 rpm

HDD - shoot out

- Seagate Desktop
- rotation speed: 7200 rpm
- average seek time: < 10 ms
- average rotation time: 4 ms

- Seagate Cheeta 15K
- rotation speed: 15000 rpm
- average seek time: < 4 ms
- average rotation time: 2 ms

HDD - shoot out

- Seagate Desktop
 - rotation speed: 7200 rpm
 - average seek time: < 10 ms
 - average rotation time: 4 ms
 - average time to read a sector: < 14 ms
- Seagate Cheeta 15K
 - rotation speed: 15000 rpm
 - average seek time: < 4 ms
 - average rotation time: 2 ms

- Seagate Desktop
 - rotation speed: 7200 rpm
 - average seek time: < 10 ms
 - average rotation time: 4 ms
 - average time to read a sector: < 14 ms
- Seagate Cheeta 15K
 - rotation speed: 15000 rpm
 - average seek time: < 4 ms
 - average rotation time: 2 ms
 - average time to read a sector: < 6 ms

HDD - shoot out

- Seagate Desktop
- rotation speed: 7200 rpm
- average seek time: < 10 ms
- average rotation time: 4 ms
- average time to read a sector: < 14 ms
- capacity: 2 TiByte

- Seagate Cheeta 15K
- rotation speed: 15000 rpm
- average seek time: < 4 ms
- average rotation time: 2 ms
- average time to read a sector: < 6 ms
- capacity: 600 GiByte

- Seagate Desktop
- rotation speed: 7200 rpm
- average seek time: < 10 ms
- average rotation time: 4 ms
- average time to read a sector: < 14 ms
- capacity: 2 TiByte
- aprx. price: 900:-

- Seagate Cheeta 15K
- rotation speed: 15000 rpm
- average seek time: < 4 ms
- average rotation time: 2 ms
- average time to read a sector: < 6 ms
- capacity: 600 GiByte
- aprx. price: 2.200:-

- Seagate Desktop
- rotation speed: 7200 rpm
- average seek time: < 10 ms
- average rotation time: 4 ms
- average time to read a sector: < 14 ms
- capacity: 2 TiByte
- aprx. price: 900:-

- Seagate Cheeta 15K
- rotation speed: 15000 rpm
- average seek time: < 4 ms
- average rotation time: 2 ms
- average time to read a sector: < 6 ms
- capacity: 600 GiByte
- aprx. price: 2.200:-
- cost capacity: 3.70 SEK/GiByte

- Seagate Desktop
- rotation speed: 7200 rpm
- average seek time: < 10 ms
- average rotation time: 4 ms
- average time to read a sector: < 14 ms
- capacity: 2 TiByte
- aprx. price: 900:-
- cost capacity: 0.44 SEK/GiByte

- Seagate Cheeta 15K
- rotation speed: 15000 rpm
- average seek time: < 4 ms
- average rotation time: 2 ms
- average time to read a sector: < 6 ms
- capacity: 600 GiByte
- aprx. price: 2.200:-
- cost capacity: 3.70 SEK/GiByte

If a sector is 512 bytes, it takes 10ms to find and read a sector, and we want to read 512 MiBytes then

If a sector is 512 bytes, it takes 10ms to find and read a sector, and we want to read 512 MiBytes then

- Time to find first sector is less relevant.

If a sector is 512 bytes, it takes 10ms to find and read a sector, and we want to read 512 MiBytes then?

- Time to find first sector is less relevant.
- If sectors that belong to the same file are close to each other we minimize movement of arm.

If a sector is 512 bytes, it takes 10ms to find and read a sector, and we want to read 512 MiBytes then?

- Time to find first sector is less relevant.
- If sectors that belong to the same file are close to each other we minimize movement of arm.
- Rotational speed should be high.

If a sector is 512 bytes, it takes 10ms to find and read a sector, and we want to read 512 MiBytes then?

- Time to find first sector is less relevant.
- If sectors that belong to the same file are close to each other we minimize movement of arm.
- Rotational speed should be high.
- The density i.e. how many sectors in each track is important.

If a sector is 512 bytes, it takes 10ms to find and read a sector, and we want to read 512 MiBytes then

- Time to find first sector is less relevant.
- If sectors that belong to the same file are close to each other we minimize movement of arm.
- Rotational speed should be high.
- The density i.e. how many sectors in each track is important.
- The communication with the drive should be fast.

If a sector is 512 bytes, it takes 10ms to find and read a sector, and we want to read 512 MiBytes then?

- Time to find first sector is less relevant.
- If sectors that belong to the same file are close to each other we minimize movement of arm.
- Rotational speed should be high.
- The density i.e. how many sectors in each track is important.
- The communication with the drive should be fast.
- Typical read and write performance is between 150 MiByte/s to 250 MiByte/s.

Historically, the Operating System was in complete control:

- it knew the layout cylinder-head-sector (CHS),

Historically, the Operating System was in complete control:

- it knew the layout cylinder-head-sector (CHS),
- could order data in segments that were close to each other and,

Historically, the Operating System was in complete control:

- it knew the layout cylinder-head-sector (CHS),
- could order data in segments that were close to each other and,
- would schedule disk operations to minimize arm movement.

Historically, the Operating System was in complete control:

- it knew the layout cylinder-head-sector (CHS),
- could order data in segments that were close to each other and,
- would schedule disk operations to minimize arm movement.

Today, the drive can often make a better decision:

Historically, the Operating System was in complete control:

- it knew the layout cylinder-head-sector (CHS),
- could order data in segments that were close to each other and,
- would schedule disk operations to minimize arm movement.

Today, the drive can often make a better decision:

- it knows, but might not reveal, the layout.

Historically, the Operating System was in complete control:

- it knew the layout cylinder-head-sector (CHS),
- could order data in segments that were close to each other and,
- would schedule disk operations to minimize arm movement.

Today, the drive can often make a better decision:

- it knows, but might not reveal, the layout.
- The operating system can help in grouping operations together, allowing the drive to decide in what order they should be done (Native Command Queuing).

Historically, the Operating System was in complete control:

- it knew the layout cylinder-head-sector (CHS),
- could order data in segments that were close to each other and,
- would schedule disk operations to minimize arm movement.

Today, the drive can often make a better decision:

- it knows, but might not reveal, the layout.
- The operating system can help in grouping operations together, allowing the drive to decide in what order they should be done (Native Command Queuing).

There is a reason why MS-DOS is called MS-DOS.

SSD - Solid State Drive

Samsung 850 EVO



- total capacity: 250 GiByte

SSD - Solid State Drive

Samsung 850 EVO



- total capacity: 250 GiByte
- form factor: 2.5"

SSD - Solid State Drive

Samsung 850 EVO



- total capacity: 250 GiByte
- form factor: 2.5"
- connection: SATA III

SSD - Solid State Drive

Samsung 850 EVO



- total capacity: 250 GiByte
- form factor: 2.5"
- connection: SATA III
- cache size: 64 MiByte

SSD - Solid State Drive

Samsung 850 EVO



- total capacity: 250 GiByte
- form factor: 2.5"
- connection: SATA III
- cache size: 64 MiByte
- random access: 30 μ s

SSD - Solid State Drive

Samsung 850 EVO



- total capacity: 250 GiByte
- form factor: 2.5"
- connection: SATA III
- cache size: 64 MiByte
- random access: 30 μ s
- read throughput: 540 MiByte/s

SSD - Solid State Drive

Samsung 850 EVO



- total capacity: 250 GiByte
- form factor: 2.5"
- connection: SATA III
- cache size: 64 MiByte
- random access: 30 μ s
- read throughput: 540 MiByte/s

aprx price, October 2018, 685:-

SD cards - flash memory

SanDisk Ultra SDXC



- form factor: SDXC

SD cards - flash memory

SanDisk Ultra SDXC



- form factor: SDXC
- capacity: 64 GiByte

SD cards - flash memory

SanDisk Ultra SDXC



- form factor: SDXC
- capacity: 64 GiByte
- read performance: 80 MiByte/s

SD cards - flash memory

SanDisk Ultra SDXC



- form factor: SDXC
- capacity: 64 GiByte
- read performance: 80 MiByte/s

aprx price, October 2016, 300:-

NAND - flash storage

memory bank




NAND - flash storage

memory bank




NAND - flash storage


memory bank




NAND - flash storage




NAND - flash storage




NAND - flash storage



NAND - flash storage



NAND - flash storage



Drive	Capacity	Price	SEK/GiByte
HDD Desktop	2 TiByte	900:-	44 öre
HDD Performance	600 GiByte	2.200:-	3.70:-
SSD Desktop	250 GiByte	685:-	2.75:-

Drive	Capacity	Price	SEK/GiByte
HDD Desktop	2 TiByte	900:-	44 öre
HDD Performance	600 GiByte	2.200:-	3.70:-
SSD Desktop	250 GiByte	685:-	2.75:-

2016 figures: SSD 4:-/GiByte

SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



- total capacity: 2 TiByte

SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



- total capacity: 2 TiByte
- form factor: 3.5"

SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm

SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA-III

SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA-III
- SSD cache: 8 GiByte

SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA-III
- SSD cache: 8 GiByte
- cache size: 64 MiByte

SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA-III
- SSD cache: 8 GiByte
- cache size: 64 MiByte
- read throughput: 210 MByte/s

SSHD - Hybrid SSD/HDD

Seagate FireCuda - SSHD



- total capacity: 2 TiByte
- form factor: 3.5"
- rotational speed: 7.200 rpm
- connection: SATA-III
- SSD cache: 8 GiByte
- cache size: 64 MiByte
- read throughput: 210 MByte/s

Seagate FireCuda SSHD, aprx price, November 2018, 1.200:-

- SATA-III - 6 Gb/s, most internal HDD and SSD today

- SATA-III - 6 Gb/s, most internal HDD and SSD today
- SAS-3 - 12 Gb/s, enterprise RAID HDD

- SATA-III - 6 Gb/s, most internal HDD and SSD today
- SAS-3 - 12 Gb/s, enterprise RAID HDD
- USB3.1 - 10 Gb/s, everything

- SATA-III - 6 Gb/s, most internal HDD and SSD today
- SAS-3 - 12 Gb/s, enterprise RAID HDD
- USB3.1 - 10 Gb/s, everything
- PCI Express 3.0 x16 - 128 Gb/s, what is it used for?

- SATA-III - 6 Gb/s, most internal HDD and SSD today
- SAS-3 - 12 Gb/s, enterprise RAID HDD
- USB3.1 - 10 Gb/s, everything
- PCI Express 3.0 x16 - 128 Gb/s, what is it used for?

An SSD has a read throughput of 500 MiByte/s which is a b/s?

SSD on the PCIe bus

Corsair Neutron NX500



- total capacity: 400 GiByte

SSD on the PCIe bus

Corsair Neutron NX500



- total capacity: 400 GiByte
- connection: PCI Express 3.0 x4

SSD on the PCIe bus


Corsair Neutron NX500



- total capacity: 400 GiByte
- connection: PCI Express 3.0 x4
- read performance: 3000 MByte/s

SSD on the PCIe bus

Corsair Neutron NX500



- total capacity: 400 GiByte
- connection: PCI Express 3.0 x4
- read performance: 3000 MByte/s
- write performance: 2400 MByte/s

SSD on the PCIe bus

Corsair Neutron NX500



- total capacity: 400 GiByte
- connection: PCI Express 3.0 x4
- read performance: 3000 MByte/s
- write performance: 2400 MByte/s

aprx price, November 2018, 3.599:-

SSD on the PCIe bus

Corsair Neutron NX500



- total capacity: 400 GiByte
- connection: PCI Express 3.0 x4
- read performance: 3000 MByte/s
- write performance: 2400 MByte/s

*aprx price, November 2018, 3.599:-
2016 October, Intel SSD 400 GB, 4.599:-*

The M.2 connector

Samsung 960 PRO 512GB



- total capacity: 512 GiByte

The M.2 connector

Samsung 960 PRO 512GB



- total capacity: 512 GiByte
- form factor: M.2-

The M.2 connector

Samsung 960 PRO 512GB



- total capacity: 512 GiByte
- form factor: M.2-
- connection: PCI Express 3.0 x4

The M.2 connector

Samsung 960 PRO 512GB



- total capacity: 512 GiByte
- form factor: M.2-
- connection: PCI Express 3.0 x4
- read performance: 3.500 MByte/s

The M.2 connector

Samsung 960 PRO 512GB



- total capacity: 512 GiByte
- form factor: M.2-
- connection: PCI Express 3.0 x4
- read performance: 3.500 MByte/s
- write performance: 2.100 MByte/s

The M.2 connector

Samsung 960 PRO 512GB




- total capacity: 512 GiByte
- form factor: M.2-
- connection: PCI Express 3.0 x4
- read performance: 3.500 MByte/s
- write performance: 2.100 MByte/s

aprx price, November 2018, 2.890:-


SSD on the memory bus

HP NVDIMM 8GB



SSD on the memory bus


HP NVDIMM 8GB



- regular DRAM backed up by Flash

SSD on the memory bus


HP NVDIMM 8GB



- regular DRAM backed up by Flash
- total capacity: 16 GiByte

SSD on the memory bus


HP NVDIMM 8GB



- regular DRAM backed up by Flash
- total capacity: 16 GiByte
- form factor: DDR4 SDIM

SSD on the memory bus

HP NVDIMM 8GB




- regular DRAM backed up by Flash
- total capacity: 16 GiByte
- form factor: DDR4 SDIM
- bus speed: 2666 MT/s


aprx price, November 2018, 7.600:-

Next year?

Intel Optane - 3D XPoint NVDIMM




Intel Optane - 3D XPoint NVDIMM




- in the pipe line

Intel Optane - 3D XPoint NVDIMM



- in the pipe line
- total capacity: 512 GiByte

Intel Optane - 3D XPoint NVDIMM



- in the pipe line
- total capacity: 512 GiByte

Increase capacity, performance and/or reliability

Redundant Array of Independent Disks RAID



Increase capacity, performance and/or reliability

Redundant Array of Independent Disks RAID



- Multiple disks that can provide:

Increase capacity, performance and/or reliability

Redundant Array of Independent Disks RAID



- Multiple disks that can provide:
- capacity: looks like a 20 TiByte disk but is actually 10 2TiByte disks

Increase capacity, performance and/or reliability

Redundant Array of Independent Disks RAID



- Multiple disks that can provide:
- capacity: looks like a 20 TiByte disk but is actually 10 2TiByte disks
- performance: spread a file across ten drives, read and write in parallel

Increase capacity, performance and/or reliability

Redundant Array of Independent Disks RAID



- Multiple disks that can provide:
- capacity: looks like a 20 TiByte disk but is actually 10 2TiByte disks
- performance: spread a file across ten drives, read and write in parallel
- reliability: write the same file to several disks, if one crashes - not a problem

Alternatives:

Alternatives:

- The cabinet that holds the disks present itself as one drive.

Alternatives:

- The cabinet that holds the disks present itself as one drive.
- A device driver in the kernel knows that we have several disks but the kernel presents it as one disk to the application layer.

Alternatives:

- The cabinet that holds the disks present itself as one drive.
- A device driver in the kernel knows that we have several disks but the kernel presents it as one disk to the application layer.
- The application layer knows that we have several disks but provides a API to other applications that looks a single drive.


- RAID 0: *stripe* files across several drives.

- RAID 0: *stripe* files across several drives.
- RAID 1: keep a complete *mirror copy* of each file.

- RAID 0: *stripe* files across several drives.
- RAID 1: keep a complete *mirror copy* of each file.
- RAID 2-6: spread a file plus parity information across several drives.

application layer, simple to understand


hardware - a complete mess



application layer, simple to understand

I/O and memory buses, protocols such as SATA, SCSI, USB etc

hardware - a complete mess




application layer, simple to understand

now it's a bit structured

I/O and memory buses, protocols such as SATA, SCSI, USB etc

hardware - a complete mess




application layer, simple to understand

device drivers that know what they are doing

now it's a bit structured

I/O and memory buses, protocols such as SATA, SCSI, USB etc

hardware - a complete mess




application layer, simple to understand

all devices have a generic API
device drivers that know what they are doing

now it's a bit structured
I/O and memory buses, protocols such as SATA, SCSI, USB etc

hardware - a complete mess



application layer, simple to understand

system calls: open, read, write, lseek ...

all devices have a generic API
device drivers that know what they are doing

now it's a bit structured
I/O and memory buses, protocols such as SATA, SCSI, USB etc

hardware - a complete mess

